

Open in app ↗

 Search WriteWriting is for everyone. [Register for Medium Day](#)

Predicting Titanic Survival



Siddharth Rao Kartik · 3 min read · Just now



A step-by-step machine learning journey using Kaggle's Titanic dataset, guided by ChatGPT, following the CRISP-DM data science framework.

Introduction

In this project, I used Kaggle's Titanic dataset to predict passenger survival. With the help of ChatGPT and the CRISP-DM framework, I explored the data, prepared it, trained machine learning models, and compared their performance. The goal was not only to build a predictive model but also to learn how AI tools like ChatGPT can guide the entire data science process.

```
import pandas as pd

df = pd.read_csv('/content/train.csv')

df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

Data Understanding

The dataset includes demographic, ticket, and travel details such as class, sex, age, fare, and cabin. Some columns contain missing values.

df.info()

df.describe(include='all').T.head(15)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
Column Non-Null Count Dtype

0 PassengerId 891 non-null int64
1 Survived 891 non-null int64
2 Pclass 891 non-null int64
3 Name 891 non-null object
4 Sex 891 non-null object
5 Age 714 non-null float64
6 SibSp 891 non-null int64
7 Parch 891 non-null int64
8 Ticket 891 non-null object
9 Fare 891 non-null float64
10 Cabin 204 non-null object
11 Embarked 889 non-null object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

	count	unique	top	freq	mean	std	min	25%	50%	75%	max
PassengerId	891.0	NaN	NaN	NaN	446.0	257.353842	1.0	223.5	446.0	668.5	891.0
Survived	891.0	NaN	NaN	NaN	0.383838	0.486592	0.0	0.0	0.0	1.0	1.0
Pclass	891.0	NaN	NaN	NaN	2.308642	0.836071	1.0	2.0	3.0	3.0	3.0
Name	891	891	Dooley, Mr. Patrick	1	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Sex	891	2	male	577	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Age	714.0	NaN	NaN	NaN	29.699118	14.526497	0.42	20.125	28.0	38.0	80.0
SibSp	891.0	NaN	NaN	NaN	0.523008	1.102743	0.0	0.0	0.0	1.0	8.0
Parch	891.0	NaN	NaN	NaN	0.381594	0.806057	0.0	0.0	0.0	0.0	6.0
Ticket	891	681	347082	7	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Fare	891.0	NaN	NaN	NaN	32.204208	49.693429	0.0	7.9104	14.4542	31.0	512.3292
Cabin	204	147	G6	4	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Embarked	889	3	S	644	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
missing = df.isna().mean().sort_values(ascending=False)
missing.head(15)
```

	0
Cabin	0.771044
Age	0.198653
Embarked	0.002245
PassengerId	0.000000
Name	0.000000
Pclass	0.000000
Survived	0.000000
Sex	0.000000
Parch	0.000000
SibSp	0.000000
Fare	0.000000
Ticket	0.000000


dtype: float64

Exploratory Data Analysis (EDA)

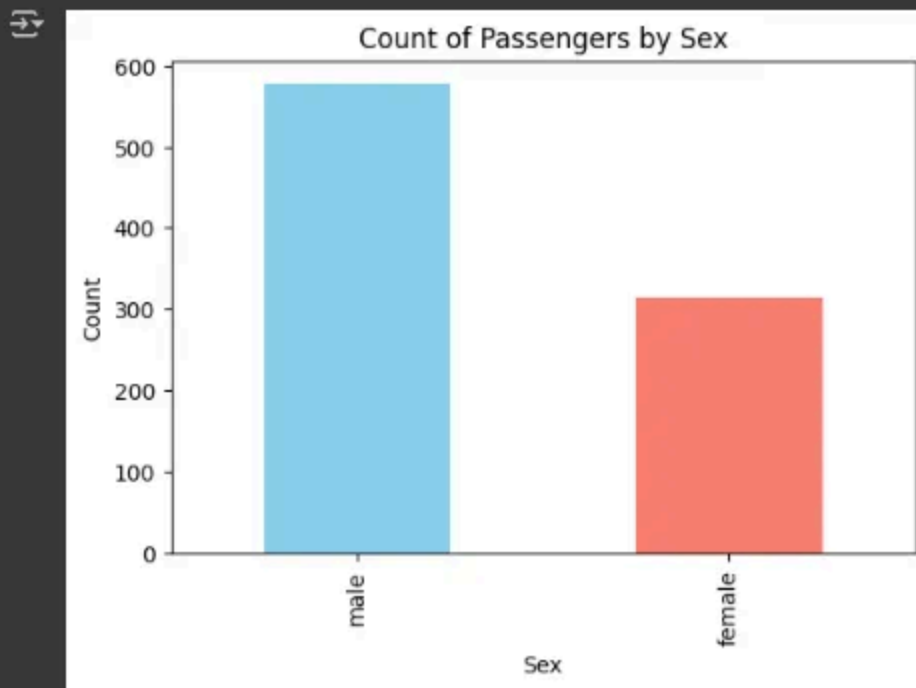
To better understand patterns, I plotted distributions:

```
import matplotlib.pyplot as plt

plt.figure(figsize=(6,4))
df['Age'].hist(bins=30, edgecolor='black')
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



```
plt.figure(figsize=(6,4))  
df['Sex'].value_counts().plot(kind='bar', color=['skyblue','salmon'])  
plt.title('Count of Passengers by Sex')  
plt.xlabel('Sex')  
plt.ylabel('Count')  
plt.show()
```



Majority of passengers were male. Age distribution is right-skewed, with many young adults.

Data Preparation

I dropped text-heavy columns (Name, Ticket, Cabin) and split the data into training and test sets (80/20).

```
from sklearn.model_selection import train_test_split

drop_cols = ['PassengerId', 'Name', 'Ticket', 'Cabin']
df_model = df.drop(columns=drop_cols, errors='ignore')

X = df_model.drop(columns=['Survived'])
y = df_model['Survived']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

len(X_train), len(X_test)
```

↔ (712, 179)

Baseline Model — Logistic Regression

I built a pipeline with imputation, scaling, and one-hot encoding, then trained a Logistic Regression model.

```

from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, f1_score, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

num_cols = X.select_dtypes(include='number').columns.tolist()
cat_cols = [c for c in X.columns if c not in num_cols]

numeric = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer([
    ('num', numeric, num_cols),
    ('cat', categorical, cat_cols)
])

clf_lr = Pipeline([
    ('preprocessor', preprocessor),
    ('model', LogisticRegression(max_iter=1000))
])

clf_lr.fit(X_train, y_train)


y_pred = clf_lr.predict(X_test)

acc = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

ConfusionMatrixDisplay.from_estimator(clf_lr, X_test, y_test)
plt.title("Confusion Matrix - Logistic Regression")
plt.show()

```

Baseline Logistic Regression -> Accuracy: 0.804 F1: 0.724



Tuned Model — Random Forest

I trained a Random Forest with GridSearchCV to tune parameters.


```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

rf_pipe = Pipeline([
    ('preprocessor', preprocessor),
    ('model', RandomForestClassifier(random_state=42))
])

param_grid = {
    'model__n_estimators': [100, 200],
    'model__max_depth': [None, 5, 10],
    'model__min_samples_split': [2, 5]
}

grid = GridSearchCV(rf_pipe, param_grid, cv=3, n_jobs=-1, scoring='accuracy')
grid.fit(X_train, y_train)

best_rf = grid.best_estimator_
y_pred_rf = best_rf.predict(X_test)

acc_rf = accuracy_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)
print("Best Params:", grid.best_params_)
print("RandomForest -> Accuracy:", round(acc_rf,3), "F1:", round(f1_rf,3))

```

Best Params: {'model__max_depth': 10, 'model__min_samples_split': 5, 'model__n_estimators': 100}
 RandomForest -> Accuracy: 0.81 F1: 0.726

Feature Importance

Random Forest highlighted the most influential predictors of survival.

```

import numpy as np

best = best_rf

ohe = best.named_steps['preprocessor'].named_transformers_['cat'].named_steps['encoder']
num_names = X.select_dtypes(include='number').columns.tolist()
cat_names = list(ohe.get_feature_names_out([c for c in X.columns if c not in num_names]))
feature_names = num_names + cat_names

importances = best.named_steps['model'].feature_importances_
top_idx = np.argsort(importances)[::-1][:10]

print("Top 10 features:")
for i in top_idx:
    print(f"{feature_names[i]:25s} {importances[i]:.4f}")

```

Top 10 features:

Fare	0.2127
Sex_male	0.1980
Age	0.1786
Sex_female	0.1738
Pclass	0.1104
SibSp	0.0454
Parch	0.0415
Embarked_S	0.0182
Embarked_C	0.0128
Embarked_Q	0.0084

Model Comparison

```
def summarize_models():  
    print("Summary:")  
    print(f"- LogisticRegression: acc={acc:.3f}, f1={f1:.3f}")  
    print(f"- RandomForest (tuned): acc={acc_rf:.3f}, f1={f1_rf:.3f}")  
    winner = "RandomForest" if acc_rf > acc else "LogisticRegression"  
    print(f"Winner: {winner}")  
  
summarize_models()
```

Summary:
- LogisticRegression: acc=0.804, f1=0.724
- RandomForest (tuned): acc=0.810, f1=0.726
Winner: RandomForest

Key Takeaways

Being female, younger, and in 1st class improved survival chances.

Fare was a strong predictor.

A tuned Random Forest outperformed Logistic Regression.

Limitations & Next Steps

Could add feature engineering (e.g., titles from names, family size)

Try advanced models like XGBoost or deep learning

Perform fairness analysis across subgroups.

Conclusion

This project showed how the CRISP-DM framework and ChatGPT can guide the entire data science process. From understanding the dataset to preparing the data, building baseline and tuned models, and extracting insights, ChatGPT helped me stay organized and efficient. Random Forest emerged as the best model, and the approach can be extended to more complex datasets.



Written by Siddharth Rao Kartik

0 followers · 1 following

Edit profile

No responses yet



Siddharth Rao Kartik

What are your thoughts?

Recommended from Medium



JW Rhoday

Professional Prompts: The Ultimate Beginner's Guide to...



Rahul Nair

From Black Box to Fair Lending: Explaining Credit Risk Models wit...